

# Advanced Machine Learning: Auto-Regressive Models

Week 3

# So Far, Discriminative Models

- Mature field
- Defining the learning task was straight forward
- Mainly served us for motivating neural architecture
- Main challenges are:
  - architecture selection
  - hyperparameter-tuning
  - But mostly, data collection and labeling

# Next: Generative Models

- Previously minor goal: put the world into buckets
- Grand goal: generating the world
- Much harder task!
- Learn a model that can generate the world
- Reconsider ImageNet
- Discriminative model: how likely is this image to contain a dog?
- Generative model: how likely is this image to be a dog image?

# Generative Modeling Tasks

- Generative models:  $P(x)$ 
  - Estimation: find density function  $Q$  that approximates  $P$
  - Sampling: draw samples from  $P$
  - Point estimation: compute the probability density of sample  $x$

# Example: ImageNet without generalization

- Estimation:  $Q(x)$  that assigns:
  - $1/N$  to every imagenet training image
  - 0 otherwise
- Sampling: sample a random ImageNet training image
- Point estimation: query  $Q$  for image  $x$

# Example: Non-Gen Results on Test Set

- Sampling: cannot sample test set images
- Point estimation: test set images give 0 probability

# Example: ImageNet with generalization

- Estimation:  $Q(x)$  - true distribution that generated ImageNet
- Sampling: from the true distribution, i.e. new ImageNet images
- Point sampling: How likely is  $x$  in the true distribution of ImageNet

# Conditional Generative Models

- Model  $P(x|y)$  rather than  $P(x)$
- Conditioning can be class labels, text prompt etc.



# Example: Class-conditioned ImageNet

- Sampling:
  - sample a new image, given its class is Dalmatian
- Point estimation:
  - how likely is this image to come from the Dalmatian dog distribution

# Why Do We Care About This?

- World models e.g., GPT (this week)
- Human-guided creativity e.g., Text-to-image (in a few weeks)

# Several Paradigms for Probability Estimation

- Auto-regressive / Non-AR (this week)
- Variational
- Adversarial
- Flow-based
- Score-based / Diffusion

# Why Learn So Many Paradigms

- All have pros and cons
- Sometimes a combination works best
- Need to choose different paradigms for situation

# Tower Rule of Probability

- $P(X, Y) = P(X|Y)P(Y)$ 
  - This is always true. Independence is when  $P(X|Y) = P(X)$
- In the case of N variables, this implies the following decomposition:

$$\begin{aligned} P(X_1, X_2, \dots, X_n) &= P(X_1)P(X_2|X_1)P(X_3|X_1, X_2) \dots \\ &= \prod_{i=1}^n P(X_i|X_1, \dots, X_{i-1}) \end{aligned}$$

# Turning This into an ML Task

- We would like to estimate  $P(x_{i+1}|x_1, x_2..x_i)$
- Assume that  $X_i$  is a categorical variable
- Learn  $Q(x_{i+1}|x_1, x_2..x_i)$
- Estimate Q via cross-entropy:

$$\ell(P(x), Q(x)) = - \sum_j P(j|x_1, x_2..x_i) \log(Q(j|x_1, x_2..x_i))$$

# Example: Language Modeling

- Input: a set of sentences decomposed into tokens
  - Here: token = word
- Estimation task: learn model that predict next token, given previous
- We **do not** observe the true  $P(j|x_1, x_2..x_i)$
- We **do** observe finite samples from  $P(j|x_1, x_2..x_i)$

## Example: Language Modeling (2)

- Sample: “the”, “cat”, “sat”, “on”, “the”, “mat”
- $Q(\text{“on”} \mid \text{“the”, “cat”, “sat”})$  – **high** likelihood
- $Q(\text{“mat”} \mid \text{“the”, “cat”, “sat”, “on”, “the”})$  – **high** likelihood
- $Q(\text{“ate”} \mid \text{“the”, “cat”, “sat”, “on”, “the”})$  – **low** likelihood



# Learning in Practice

- Given a set  $S$  containing sequences  $((x_1, x_2..x_i), y)$
- Learn classifier:

$$L = - \sum_{(x_1, x_2..x_i), y \in S} \sum_j 1_{y=j} \log(Q(j|x_1, x_2..x_i))$$

- Or more simply:

$$L = - \sum_{(x_1, x_2..x_i), y \in S} \log(Q(y|x_1, x_2..x_i))$$

# Language Modeling (2)

- In our example, we had a large set of (sentence, next word)
- Define each word as a class
- Learn a classifier that takes a sequence of words and returns class
- It should handle different sequence lengths  $0 < i \leq N$

$$L = - \sum_{(x_1, x_2 \dots x_i), y \in S} \log(Q(y|x_1, x_2 \dots x_i))$$

# Running Example

- The training sample “the cat sat on the mat”
- Transform to  $Q(y | \text{“the”, “cat”, “sat”, “on”, “the”})$
- Many possible classes: “the”, “cat”, “sat”, “on”, “ate”, “apple”, “dog” etc.
- Objective: loss cross entropy with  $(0, 1, 0, 0, 0, 0, 0, 0, \dots)$

$$L = - \sum_{(x_1, x_2 \dots x_i), y \in S} \sum_j 1_{y=j} \log(Q(j | x_1, x_2 \dots x_i))$$

# Sampling from an AR Model

- Assume we trained an AR model  $Q(j|x_1, x_2..x_i)$
- We want to generate a new token sequence  $x_1..x_N$
- Objective: generate the **most likely** sequences
- Option 1: select the most likely token
  - Lacks diversity

# Sampling New Data (2)

- To sample multiple likely sequence, we need stochasticity
- Option 2: sample according to the distribution of  $Q(j|x_1, x_2..x_i)$ 
  - high diversity, but maybe too much
- Option 3: sample from the top-K tokens
  - K can tune the amount of diversity

# Point Estimation

- Provided a model, point estimation is quite easy with AR models
- The point probability is simply:

$$Q(X) = \prod_{i=1}^N P(x_i | x_1, x_2 \dots x_{i-1})$$

- Evaluating log probability is more numerically stable

# Extension to Conditional AR Models

- The framework is easily extended to conditional probs.
- The tower rule in this case:

$$Q(X|c) = \prod_{i=1}^N P(x_i|x_1, x_2 \dots x_{i-1}, c)$$

- Model estimation, sampling, point estimation: virtually unchanged

# Perplexity: Measure of Model Quality

- Test perplexity: the geometric mean of Q on the test set

$$\prod_{m=1}^M Q(X^m)^{-\frac{1}{M}}$$

- Intuition: samples of real data should have high likelihood
- In our example:
  - Test sentences: high probability
  - Garbled sentences: low probability



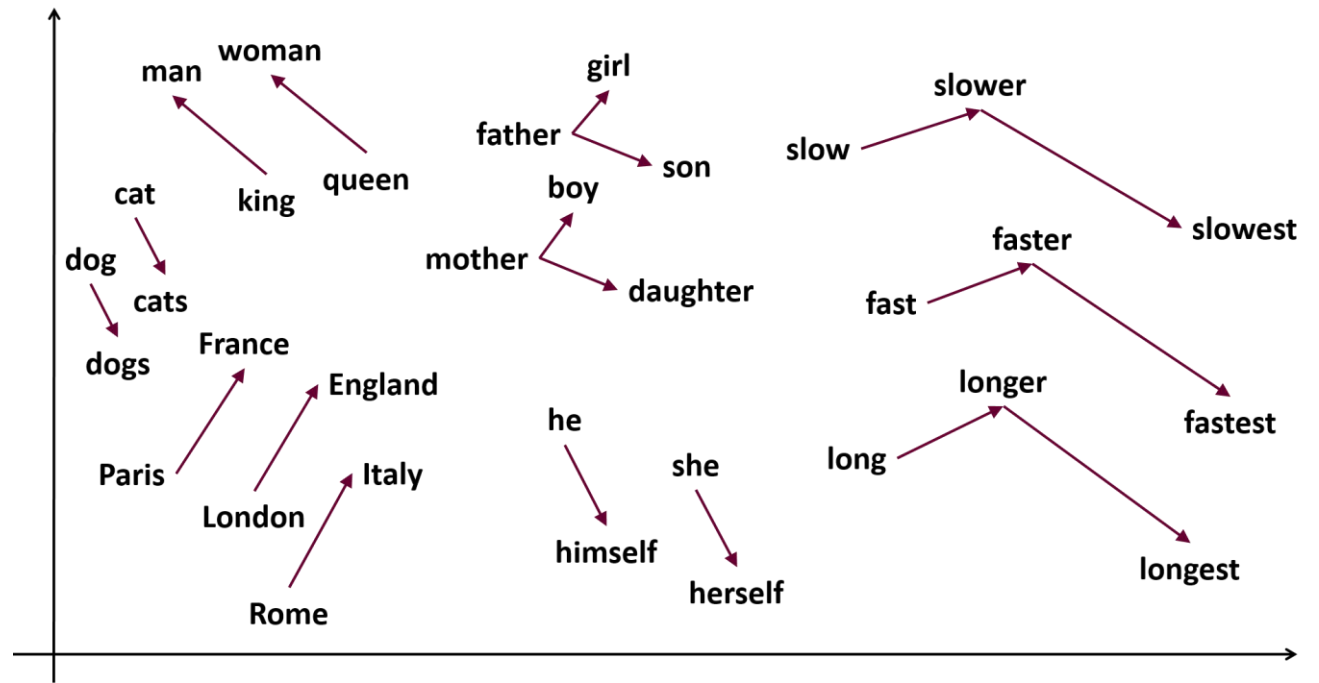
# Simple Example: Linear Language Model

- Every token  $t$  is described by a one-hot vector  $\mathbf{e}_t$
- A linear layer  $\mathbf{W}$  transforms each token into a dense vector  $\mathbf{W}\mathbf{e}_t$
- Features of all tokens are average pooled into a fixed representation
- The average vector is linearly classified to the next word matrix  $\mathbf{V}$

$$\text{softmax}(V^T W \sum_{t=1}^N \frac{e_t}{N})$$


# Intuition: Word Representations

- Matrix  $W$  embeds words into a common vector space
- Matrix  $V$  maps representations into words
- The columns of  $W$  correspond to the embeddings of each word
- These can be semantic



# Transformer Language Models

- Modern language models use huge transformers
- Much more expressive than the simple linear model
- Can look at huge context sizes (next word given last 4000 words)
- Take word order into account




# GPT3 (OpenAI)

- Similar to what is used in ChatGPT
- Context: 2048 tokens
- Parameter number: 175B
- Transformer block number: 96
- Training data: 0.5 Tn words crawled from the internet



# Are LMs the key to Sentient Machines?

- "Once LMs can predict all sentences, we can ask them virtually any question and get the correct answer, making them practically sentient" – what do you think?



# Conditional LMs: Machine Translation

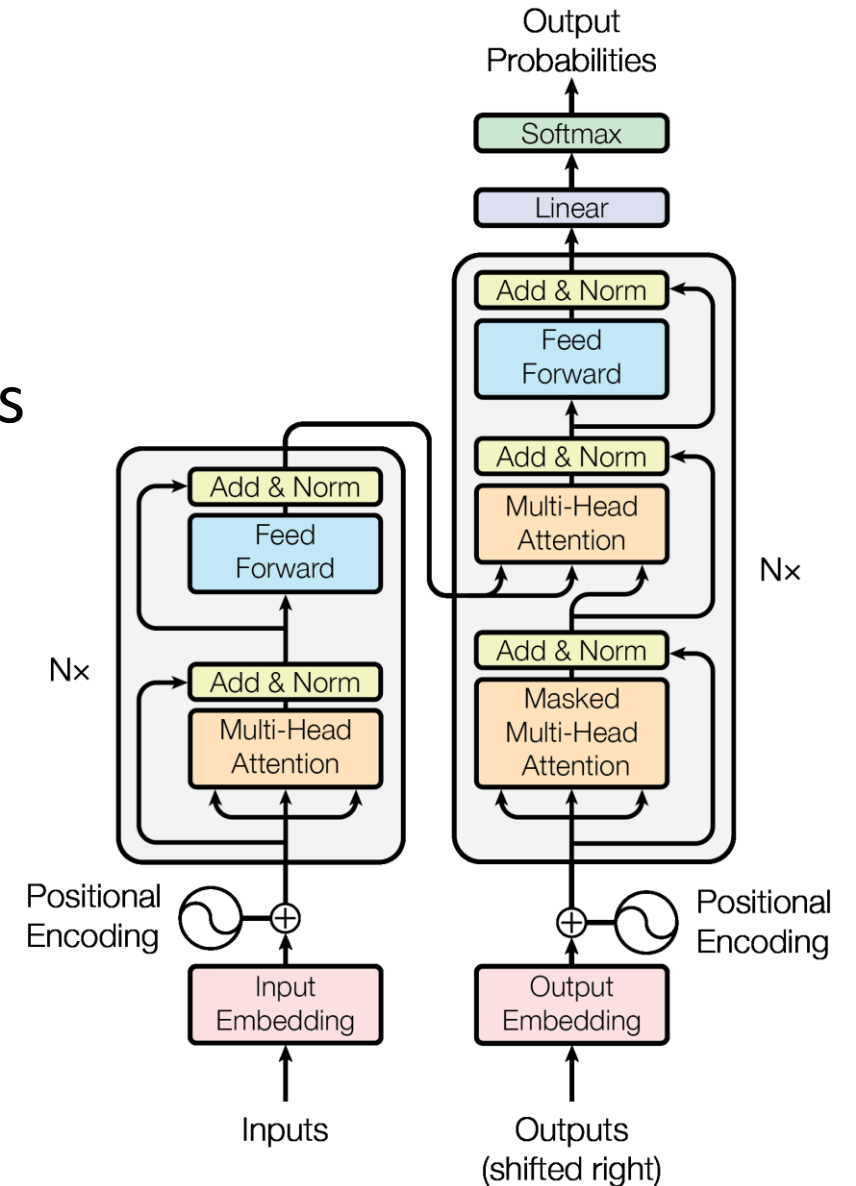
- Predict next word given previous words + source language sentence

$$Q(X|c) = \prod_{i=1}^N P(x_i|x_1, x_2..x_{i-1}, c)$$

- About the same solution as before
- The paper that started transformers:
  - “Attention is all you need”, Google 2017

# Transformer Implementation

- Transformer encoder of source sentence
- Transformer decoder over encoding + previous words in target sentence
- Why not a decoder over source sentence + previous target sentence words?



# Applying the Same Idea to Continuous Data

- LM presented so far apply to discrete multivariate data
- How can we apply this to continuous data? E.g. images
- Idea: quantization
  - Learn encoder that maps image to a small number of discrete variables
  - Need to also have a decoder that maps discrete tokens back to image
  - Then everything else applies as before



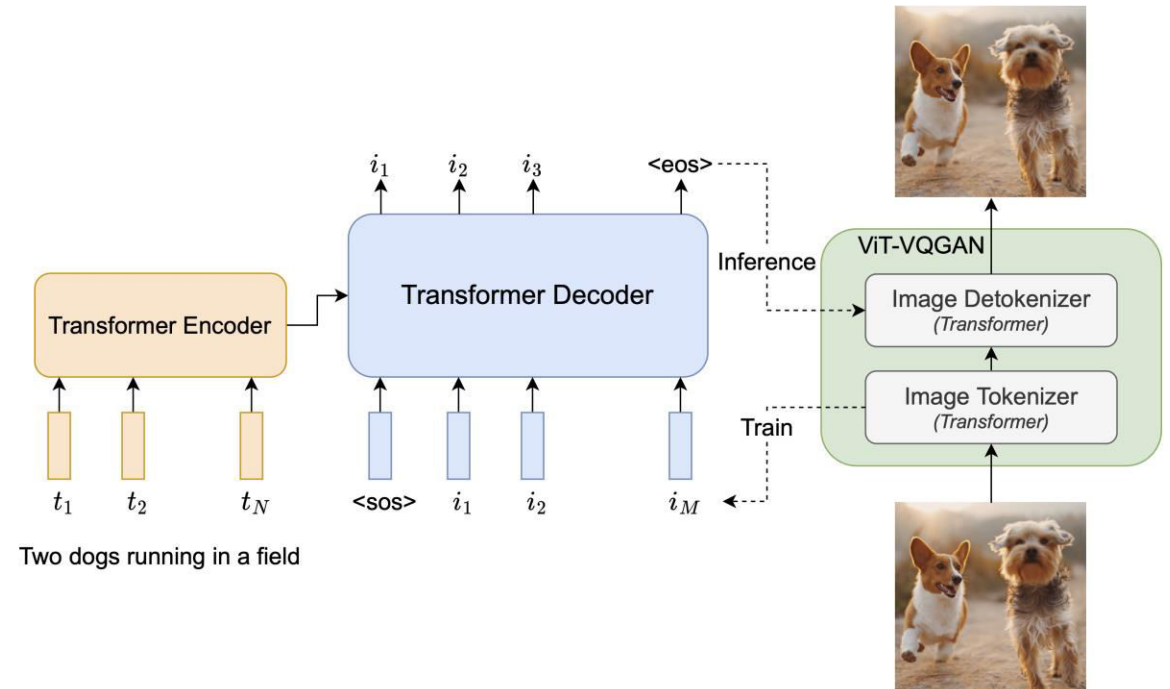
# Parti: LM for Text-to-Image Translation



A photo of an Athenian vase with a painting of *pandas* *toucans* *pangolins* playing *tennis* *soccer* *basketball* in the style of *Egyptian hieroglyphics*.

Parti

## Parti (Google)



# Effect of Number of Parameters

- Current models have poor sample and compute complexity

**350M**



**750M**



**3B**



**20B**



*A portrait photo of a kangaroo wearing an orange hoodie and blue sunglasses standing on the grass in front of the Sydney Opera House holding a sign on the chest that says Welcome Friends!*

# Code Example

- <https://github.com/karpathy/minGPT>

# Neural Scaling Laws

- Imagine you are a manager in a large company developing LMs
- You want: best LM in the world
- Have resources but they are finite
- Where would you invest them?
  - Architecture and optimization research?
  - More data?
  - More compute?
  - More memory?

# Power Law

- Power law: simple way of representating a relation between variable
- Product of powers of the variables
- Scale invariant

$$f(cx) = a(cx)^{-k} = c^{-k} f(x) \propto f(x),$$

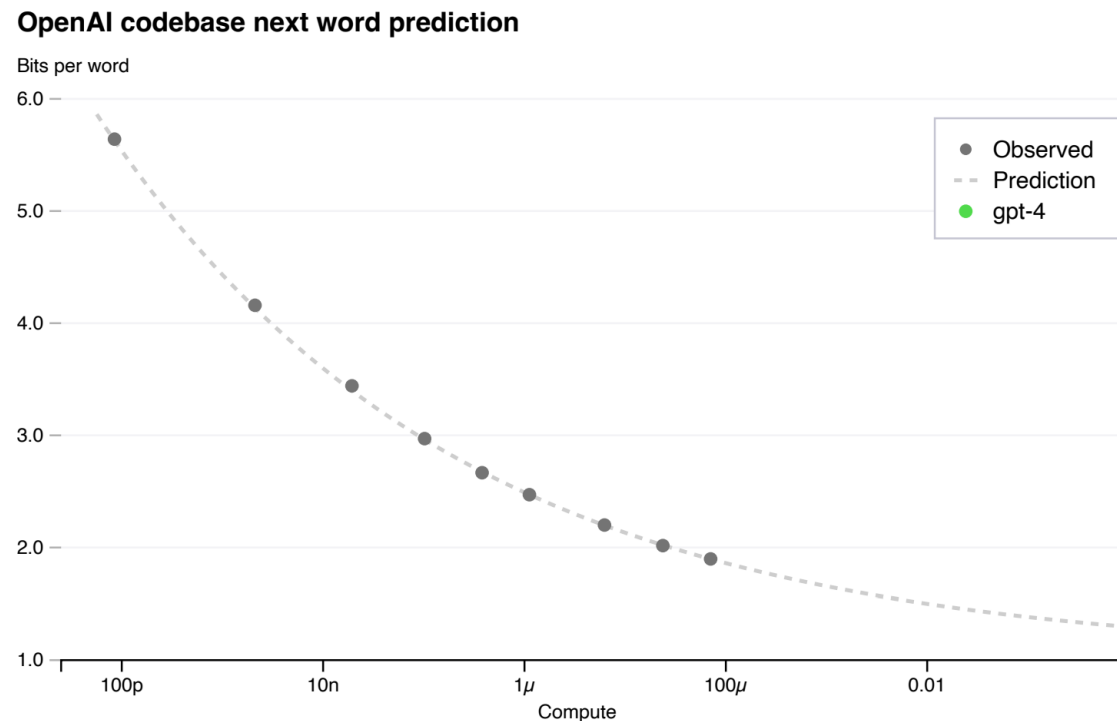
# Example: GPT4

- How can we know the accuracy of GPT4 on a small budget?
- Idea: train a set of smaller models, estimate the power law

$$L(C) = aC^b + c,$$


# GPT4 Results

- Power law predicted true accuracy very accurately
- Prediction extrapolating 10000X



# Other Findings in Previous Papers

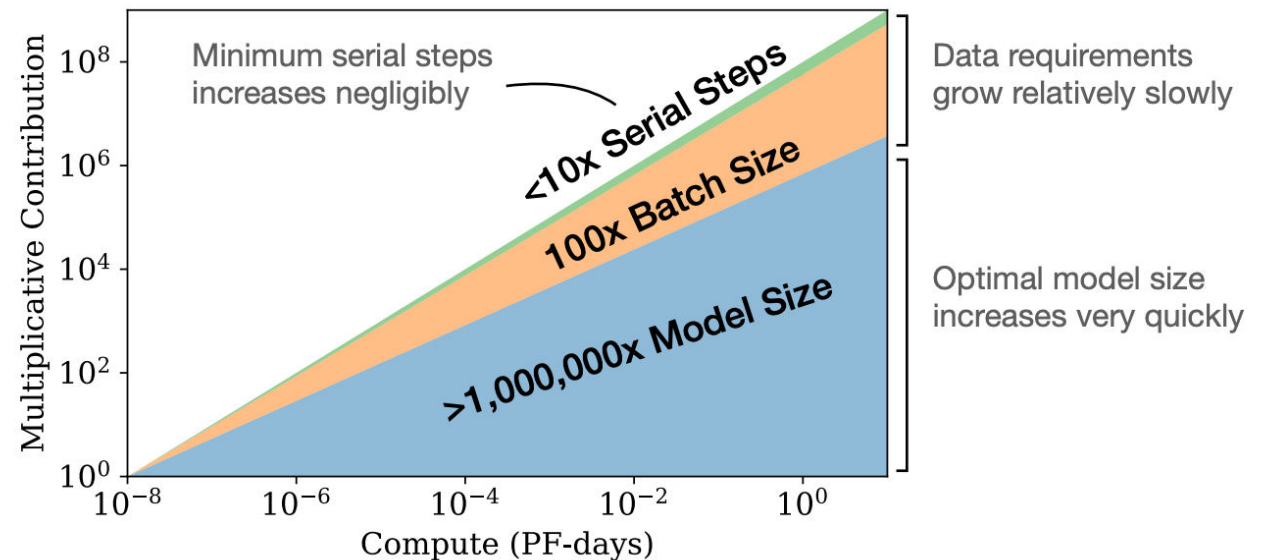
- Only total compute matters – doesn't matter if it is in depth or width
- Larger models require fewer training step to reach given loss





# What Should You Invest In?

- Given a compute budget:
  - Scale model a lot
  - Increase dataset size a little
  - Moderately increase batch



# Conclusion

- Language models are a very simple but effective density estimator
- Can be conditional or unconditional
- Given massive compute, can achieve amazing results Not suitable for all types of data
- Scaling laws predict their performance very well